

CS5800: B+-Trees

Ricardo Baeza-Yates & Anurag Bhardwaj

Northeastern University 2018

CS 5800, B-Trees, Baeza-Yates & Bhardwaj, NEU Silicon Valley –

Agenda

- B+-Trees
 - Introduction
 - Searching, Insertion, Deletion
 - Complexity Analysis
- B+-Tree Advantages

how memory works? Secondary, primary (OS) disk?
purpose of RAM? --> operating system load from RAM so it's fast
database to query efficiently? use B+ trees to solve disk storage problem

CS 5800, B-Trees, Baeza-Yates & Bhardwaj, NEU Silicon Valley2–

B+-Trees - Introduction

- Disk Storage Problem
 - How to store data on disk for faster access?
 - How about Binary Search trees?
 - # of reads needed?
 - # of root -> leaf traversals?
 - Locality of reference?
- B-Trees and its variants comes to rescue
- B-Tree is a variant of search tree with many children

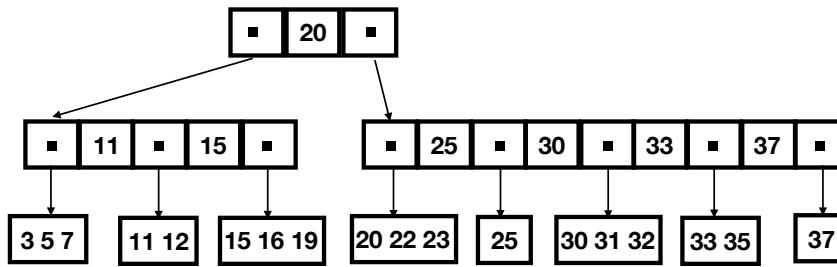
CS 5800, B-Trees, Baeza-Yates & Bhardwaj, NEU Silicon Valley3–

B+-Trees - Properties

- Key Ideas
 - Reduce # of reads: store more data in each node
 - Reduce # of traversals: high branching factor ($B \gg 2$)
 - Locality of reference: sorted keys

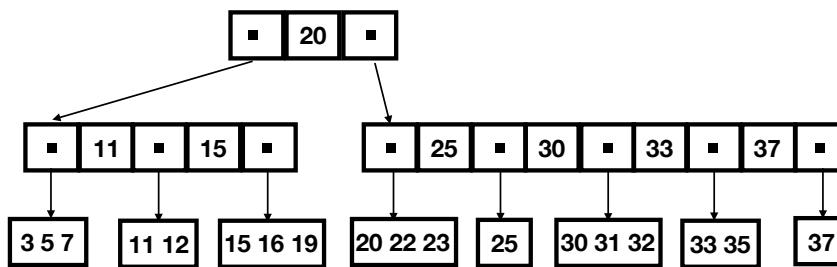
CS 5800, B-Trees, Baeza-Yates & Bhardwaj, NEU Silicon Valley4–

B+-Trees - Example



CS 5800, B-Trees, Baeza-Yates & Bhardwaj, NEU Silicon Valley –

B+-Trees - Search



- Linear or binary search in each node - given max. # of keys in node
- If not found, follow the appropriate child pointer
- Complexity: $\log_B(N)$

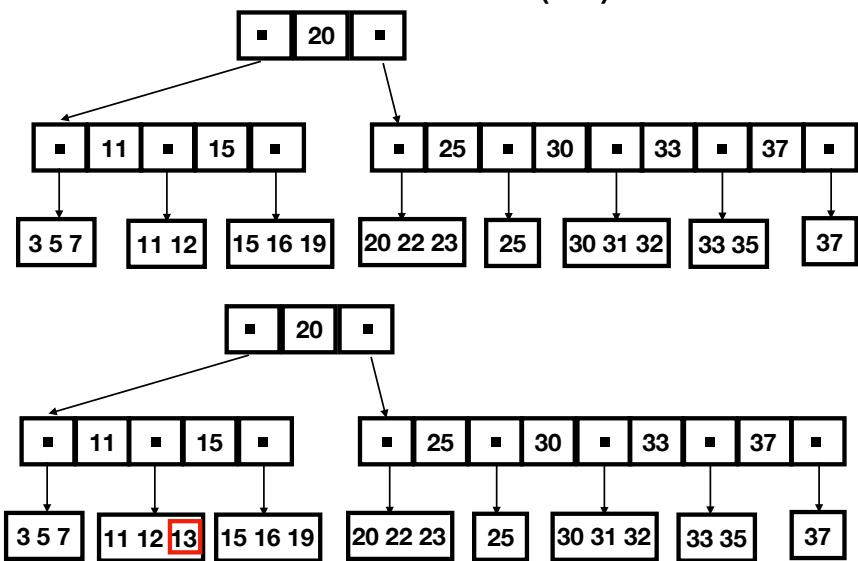
CS 5800, B-Trees, Baeza-Yates & Bhardwaj, NEU Silicon Valley –

B+-Trees - Formal Definition

- For non-leaf node N:
 - (# of children in K) - (# of keys stored in K) == 1
- For every non-root node N: store sorted array in each node
 - at least $(B-1)$ keys
- All nodes have at most $2*B$ children
- All leaf nodes have same depth
- Only leaf node(s) contain data, non-leaf contain routing values

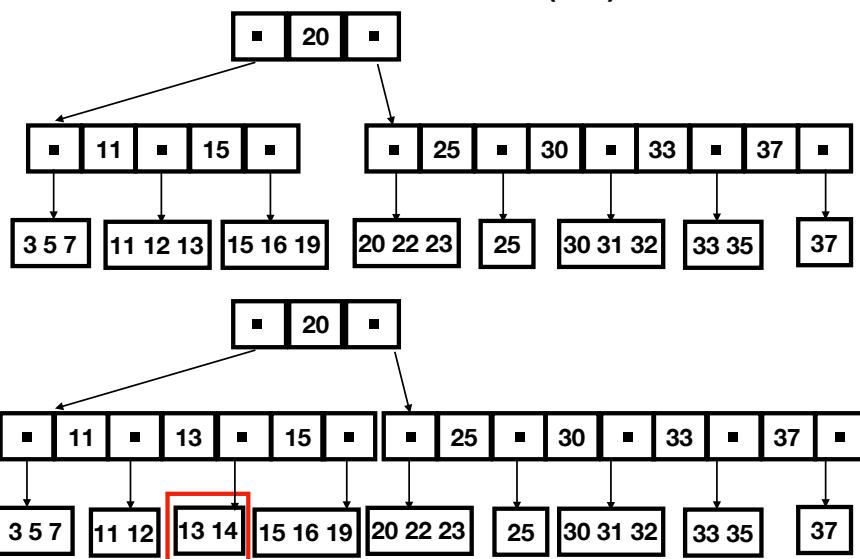
CS 5800, B-Trees, Baeza-Yates & Bhardwaj, NEU Silicon Valley –

B+-Trees - Insertion (1/3)



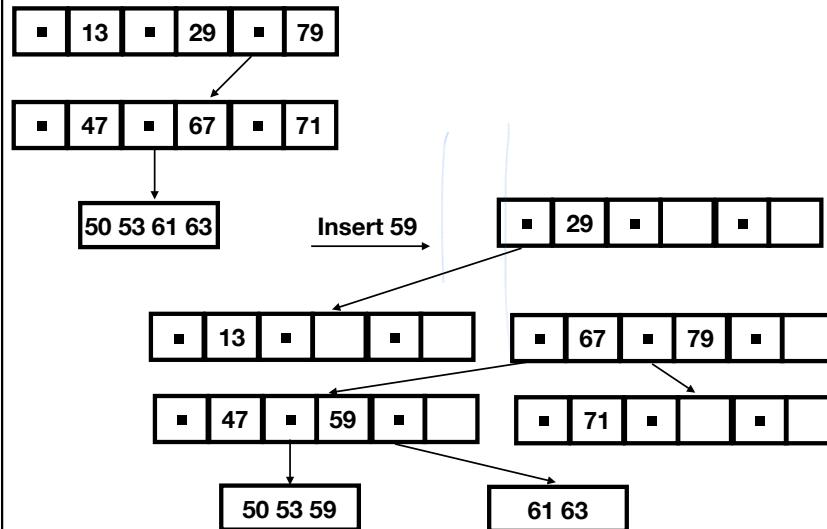
CS 5800, B-Trees, Baeza-Yates & Bhardwaj, NEU Silicon Valley –

B+-Trees - Insertion (2/3)



CS 5800, B-Trees, Baeza-Yates & Bhardwaj, NEU Silicon Valley 9

B+-Trees - Insertion (3/3)



CS 5800, B-Trees, Baeza-Yates & Bhardwaj, NEU Silicon Valley – 10

B+-Trees - Deletion

- An element is removed from the leaf
 - if leaf becomes empty, remove key from parent
 - Keep adjusting parent and its siblings to satisfy invariants
 - May ripple to root and reduce height of tree by 1
 - If the node has < 50% occupancy, merge it with left/right node
(standard B+-Trees have a minimal occupancy of 50%)

Node has maximum b keys
Each node has maximum b-2 keys
Otherwise have to merge

- More examples:
 - <https://www.cs.usfca.edu/~galles/visualization/BTree.html>

CS 5800, B-Trees, Baeza-Yates & Bhardwaj, NEU Silicon Valley 1-

B+-Trees - Complexity

- Time Complexity

B is the max number of children can have

 - Search: $O(B * \log_B(n))$
 - Insertion: $O(B * \log_B(n))$
 - Deletion: $O(B * \log_B(n))$

Space Complexity

- $O(n)$

CS 5800, B-Trees, Baeza-Yates & Bhardwaj, NEU Silicon Valley 2-

B+-Trees - Advantages

- Multi-level Indexing
- Bulk Insertion
- Aggregation queries
- Value of B can be set to match the size of a disk sector or its multiples

give an example of an application when B+ tree is more useful than binary search tree

tell me if this is a valid B+ tree or not

number of keys / children have for B+ tree?